

MagicBox Handoff

Handoff by Stanford Code the Change of Magicbox
App and Mobility Pipeline to UNICEF

2019-05-30

Copyright © 2019 Members of the 2018-2019 Stanford Code the Change UNICEF Team

Magicbox App

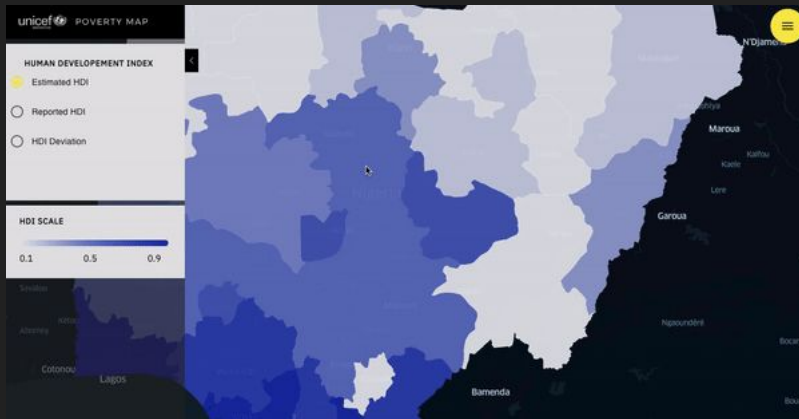
<https://github.com/codethechange/magicbox-app>

Purpose

Change the map colors when clicking on provinces

Generate layer information on mobility for province clicked

What we have



Where to go from here

Update the rest of the data store when province clicked

Condition click on specific province selected

Generate layers dynamically from CSV files

Suggestions

For UNICEF

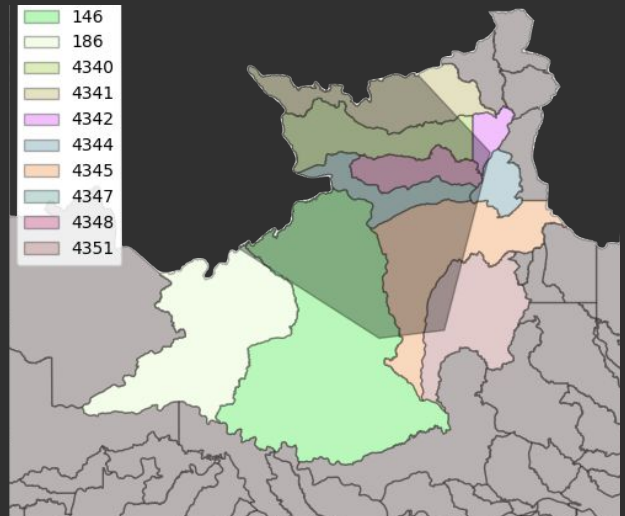
- More documentation and comments on code

For us

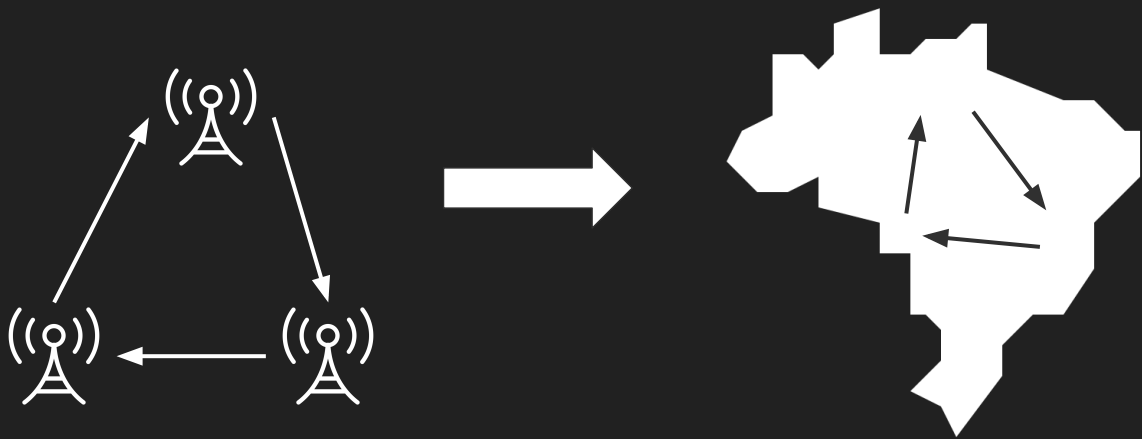
- Better understanding of Redux

Mobility Pipeline

github.com/codethechange/mobility_pipeline



Purpose



tower cell by Maxim Kulikov from the Noun Project, Brazil by Hea Poh Lin from the Noun Project

You can get data on how cell phones move between cell tower coverage areas. We can assume that this represents people moving between those areas, but we would like to know how people are moving between administrative regions like provinces within a country. In other words, we want to translate the data into the language of political regions.

The overall strategy we use to do this is detailed here in Mike's article:

<https://medium.com/@mikefabrikant/cell-towers-chiefdoms-and-anonymized-call-detail-records-a-guide-to-creating-a-mobility-matrix-d2d5c1bafb68>

Goals



fast by Andrejs Kirma from the Noun Project, Clean by Rudez Studio from the Noun Project, convenient by John Francis T from the Noun Project

We wanted to make sure our code was reasonably fast, even on large countries like Brazil. We also tried to keep our code clean and clear so that it would be easy to maintain. Finally, we wanted to keep the user experience as simple and easy as possible.

Current State



Right now, our code successfully transforms the mobility data. It runs reasonably quickly (2 mins for Brazil level 2) because of our reliance on matrix operations and a geospatial data structure called an RTree. Our code is, I think, pretty clean and readable, and we have written comprehensive tests and documentation to facilitate easy maintenance. The user experience is also simple--just run the script!

How Our Code Works

Overview of Mobility Pipeline

- **Function:** Creates the tower-admin matrix, tower-tower matrix, and admin-tower matrix, and multiplies them together (in the respective order)
- **Input:** JSON Shape Files for Administrative Region (Brazil), Voronoi Shape Files, Mobility Data (provided by UNICEF)
- **Output:** The admin-admin matrix, depicting movement of people from one admin region to another, in a `.csv` file

Computing the Tower-to-Tower Matrix

From	To	# ppl
0	0	105
0	1	53
1	1	76

From	To	# ppl
0	0	
0	1	
1	0	
1	1	

From	To	# ppl
0	0	105
0	1	53
1	0	0
1	1	76

	From		
To			
		105	0
		53	76

We get mobility data in this format, where each row is for a unique pair of towers. The rows are ordered in from-major order, so we first see all the pairs with a from tower 0, and then those with from tower 1. All tower numbers are in strictly increasing order. You can see how we might be able to re-shape the right-most row into the tower-to-tower matrix on the far right, but there's a problem because any rows with a number of people equal to 0 are missing.

Computing the Tower-to-Tower Matrix

From	To	# ppl
0	0	105
0	1	53
1	1	76

From	To	# ppl
0	0	
0	1	
1	0	
1	1	

From	To	# ppl
0	0	105
0	1	53
1	0	0
1	1	76

	From		
To			
		105	0
		53	76

To fix this problem, we create a matrix with every row, but no mobility data. We can generate this very efficiently with Pandas.

Computing the Tower-to-Tower Matrix

From	To	# ppl
0	0	105
0	1	53
1	1	76

From	To	# ppl
0	0	
0	1	
1	0	
1	1	

From	To	# ppl
0	0	105
0	1	53
1	0	0
1	1	76

	From		
To			
		105	0
		53	76

Then, we can efficiently merge these two matrices using pandas. After filling in the empty cells with zeroes, we get a matrix with all our mobility data and no missing rows.

Computing the Tower-to-Tower Matrix

From	To	# ppl
0	0	105
0	1	53
1	1	76

From	To	# ppl
0	0	
0	1	
1	0	
1	1	

From	To	# ppl
0	0	105
0	1	53
1	0	0
1	1	76

To	From		
		0	1
	0	105	0
	1	53	76

Finally, we can reshape the number of people column into a square matrix that shows all our mobility data. This is the tower-to-tower matrix. Notice how the “to” and “from” tower indices are just the row and column indices of each cell.

Tower-Admin Matrix

RTree Algorithm: determines whether two rectangles overlap

Tower-Admin Matrix:

- (1) Load RTree with list of admin region polygons and their bounding boxes
- (2) Input rectangular bounding box of Voronoi cell into RTree
- (3) Output a list of polygons whose bounding boxes overlap with that of the Voronoi
- (4) (Hash Map) ...
- (5) Compute the percentage of overlap between each polygon and Voronoi cell
- (6) Repeat (2)-(5) for all Voronoi cells

Recovering Polygon Indices

```
dict = {} # dict = {polygon_key: index}

for every admin MultiPolygon mpol with index i:

    polygon_key = Tuple(

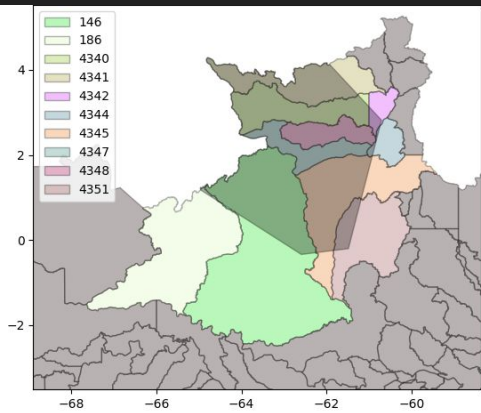
        [ Tuple(pol.exterior.coords) for pol in mpol ]

    )

    dict[polygonKey] = i
```

One complication is that the RTree returns the polygon itself, so we need a way to get its index. As an aside, we were actually considering using another library that would give us the indices, but that required installing non-Python dependencies, and we wanted to keep things simple. Anyway, we created a dictionary from polygons back to indices. Polygons aren't hashable, though, so we couldn't use them as keys. Instead, we took the coordinates of the exterior edge of the polygon and made them into a tuple. We did this for every polygon in the MultiPolygon, and but all of those tuples into another tuple. Tuples are hashable, so we use them as the key.

Tower-Admin Matrices



Region	Percent Overlap
146	0.3438314729373602
186	0.005168094316897268
4340	0.9791906016675118
4341	0.7525484088364074
4342	0.24378993279982075
4344	0.07680992472243754
4345	0.6068501217581207
4347	1.0
4348	1.0
4351	0.14599840959396204

Admin-Tower Matrix

Constructed using a similar algorithm as the Tower-Admin matrix (namely, the RTree), except the input/outputs are “flipped”:

- Load RTree with a list of Voronoi cells and their bounding boxes
- Input the bounding box of an admin region and find the percentage overlap with each Voronoi cell

Code Organization

Code Structure

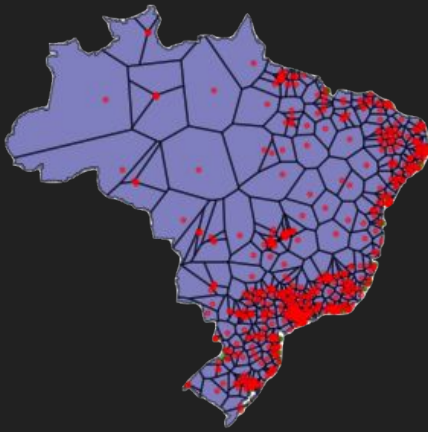
```
mobility_pipeline/  
  data/actual/  
    br_admin2.json, ...  
  docs/  
    *.rst documentation files for Sphinx  
    _build/  
      *.html generated documentation files from Sphinx  
mobility_pipeline/  
  lib/  
    make_matrix.py, ...  
    data_interface.py  
    mobility_matrix.py, ...  
  tests/src/  
    test_make_matrix.py, ...
```

In our repository, you can see at the top here our data files and documentation. The orange files here don't have any code. Then we have library files like `make_matrix`. The code in these files doesn't care about the format of the data files, and they mostly have unit tests. All the code dealing with the format of the data files lives in `data_interface`. Then we have the script files, in green, which load the data using the data interface and perform operations using the library files. The idea here is that the library files could be re-used in other contexts, while the rest of the code is tightly linked to this particular application. Finally we have the test files, with one test file for each library file.

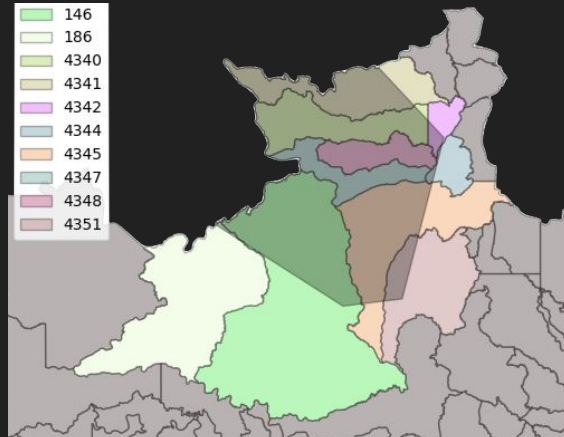
Everything Else Included

Visualization Tools

Note: For `matplotlib` to work on macOS, you have to create the virtual environment with `python -m venv` instead of `virtualenv`.



```
python plot_voronoi.py
```



```
python visualize_overlaps.py
```

We created some visualization tools along the way to sanity-check our work, and we thought you might find them useful for seeing what the data really mean. `Plot_voronoi` shows the Voronoi tessellation with the towers overlaid.

`visualize_overlaps` shows part of the computation behind the tower-to-admin matrix. It shows which admin regions might overlap with a Voronoi tessellation (the gray shaded region), and it prints out the matrix values computed for each admin region. The number labels are indices, which are also printed alongside the matrix values.

Finally, we created a script to validate data files. This script is slow, but it runs a variety of check to ensure that the data files look reasonable.

Documentation with Sphinx

<https://unicef-mobility-pipeline.readthedocs.io/>

The image displays two side-by-side screenshots of the UNICEF Mobility Pipeline documentation website, which is built using Sphinx and hosted on Read the Docs.

Left Screenshot: 'Using mobility_pipeline' page

- Header:** UNICEF Mobility Pipeline, latest version.
- Search:** Search docs.
- Contents:** Getting Started, Contributing, User Manual, Setup, Running the Program, Running Utilities, Developer Manual, Code Documentation.
- Section: Using mobility_pipeline**
 - Setup**
 - Getting Shapefiles**

Download the shapefile (.shp) from GADM (Database of Global Administrative Areas). After unpacking the zip file, be sure to choose the .shp file that corresponds to the level of admin you want (e.g. states versus counties in the United States).

Next, convert the .shp file into a GeoJSON format via the command line tool ogr2ogr, available by installing GDAL. Then, do the conversion with this command:

```
ogr2ogr -f GeoJSON -t_srs EPS:84 [name].geojson [name]
```

replacing [name] to match your .shp file.
 - Configuration**

Regardless of what you want to use mobility_pipeline for, you will need to tell it how to find your data. You can do so by adjusting the constants in mobility_pipeline.data_interface. These constants are:

[Read the Docs](#) [v: latest](#)

Right Screenshot: 'mobility_pipeline.visualize_overlaps module' page

- Header:** UNICEF Mobility Pipeline, latest version.
- Search:** Search docs.
- Contents:** Getting Started, Contributing, User Manual, Developer Manual, Code Documentation.
- Section: mobility_pipeline.visualize_overlaps module**

Plots one Voronoi cell on top of all the administrative regions. The admins that intersect the Voronoi cell are colored by index and have the associated values from the tower-to-admin matrix printed. This lets you check that the matrix values seem reasonable.
- mobility_pipeline.visualize_overlaps.I_TOWER_TO_COLOR**

Index of Voronoi cell to show.
- mobility_pipeline.visualize_overlaps.main()**

Main function that generates the plot
- mobility_pipeline.visualize_overlaps.plot_polygon(axes: matplotlib.pyplot.axes, polygon: shapely.geometry.multipolygon.MultiPolygon, color: str, _label: str) → None**

Plot a polygon (or multipolygon) with matplotlib

Parameters

 - axes** – The matplotlib axes to plot on
 - polygon** – The polygon to plot
 - color** – Color to use for shading the polygon
 - _label** – Label for the polygon that will be displayed in the legend

Returns

None

[Read the Docs](#) [v: latest](#)

We have detailed code comments for all functions, modules, and constants. We also wrote introductory documentation for both users and developers. These forms of documentation are combined by Sphinx, which can be built into local HTML. It is also hosted online by readthedocs.

Unit Tests with PyTest



./test.sh

Test by rajakumara from Noun Project

```
$ python -m pytest
===== test session starts =====
platform linux -- Python 3.7.1, pytest-4.4.0, py-1.8.0, pluggy-0.9.0
hypothesis profile 'default' -> database=DirectoryBasedExampleDatabase('/home/travis/build/codethechango/mobility_pipeline/
/hypothesis/examples')
rootdir: /home/travis/build/codethechango/mobility_pipeline, inifile: pytest.ini
plugins: cov-2.6.1, hypothesis-4.16.0
collected 23 items

tests/src/test_make_matrix.py ..... [ 20%]
tests/src/test_overlap.py ..... [ 30%]
tests/src/test_validate.py ..... [ 60%]
tests/src/test_voronoi.py ..... [ 90%]

----- coverage: platform linux, python 3.7.1-final-0 -----
Name                               Stmts   Miss  Cover
-----
mobility_pipeline/lib/__init__.py      0      0   100%
mobility_pipeline/lib/make_matrix.py   35      0   100%
mobility_pipeline/lib/overlap.py        5      0   100%
mobility_pipeline/lib/validate.py     97     46    54%
mobility_pipeline/lib/voronoi.py       19      1    95%
TOTAL                                157     46    73%

===== 23 passed in 5.54 seconds =====
The command "python -m pytest" exited with 0.

$ find mobility_pipeline -name "*.py" | xargs pylint
Your code has been rated at 10.00/10

The command "find mobility_pipeline -name "*.py" | xargs pylint" exited with 0.

$ find tests/src -name "*.py" | xargs pylint
Your code has been rated at 10.00/10

The command "find tests/src -name "*.py" | xargs pylint" exited with 0.

$ python -m mypy mobility_pipeline
The command "python -m mypy mobility_pipeline" exited with 0.

$ licheck -s strategy.ini -r requirements.txt
gathering licenses...
64 packages and dependencies.
check authorized packages...
64 packages.
The command "licheck -s strategy.ini -r requirements.txt" exited with 0.

$ bash <(curl -s https://codcov.io/bash)
Done. Your build exited with 0.
```

Excluding the data validation code, data interface, and scripts, we have unit tests for almost all of our codebase (98%). We don't comprehensively test edge cases though. Executing the test.sh script runs the unit tests, checks style with a linter, uses type annotations to type-check our code, and checks license compatibility. These tests are run by Travis-CI, which gives output like that shown on the right.